

Part III

Synchronization

Software and Hardware Solutions

Computers are useless. They can only give answers.

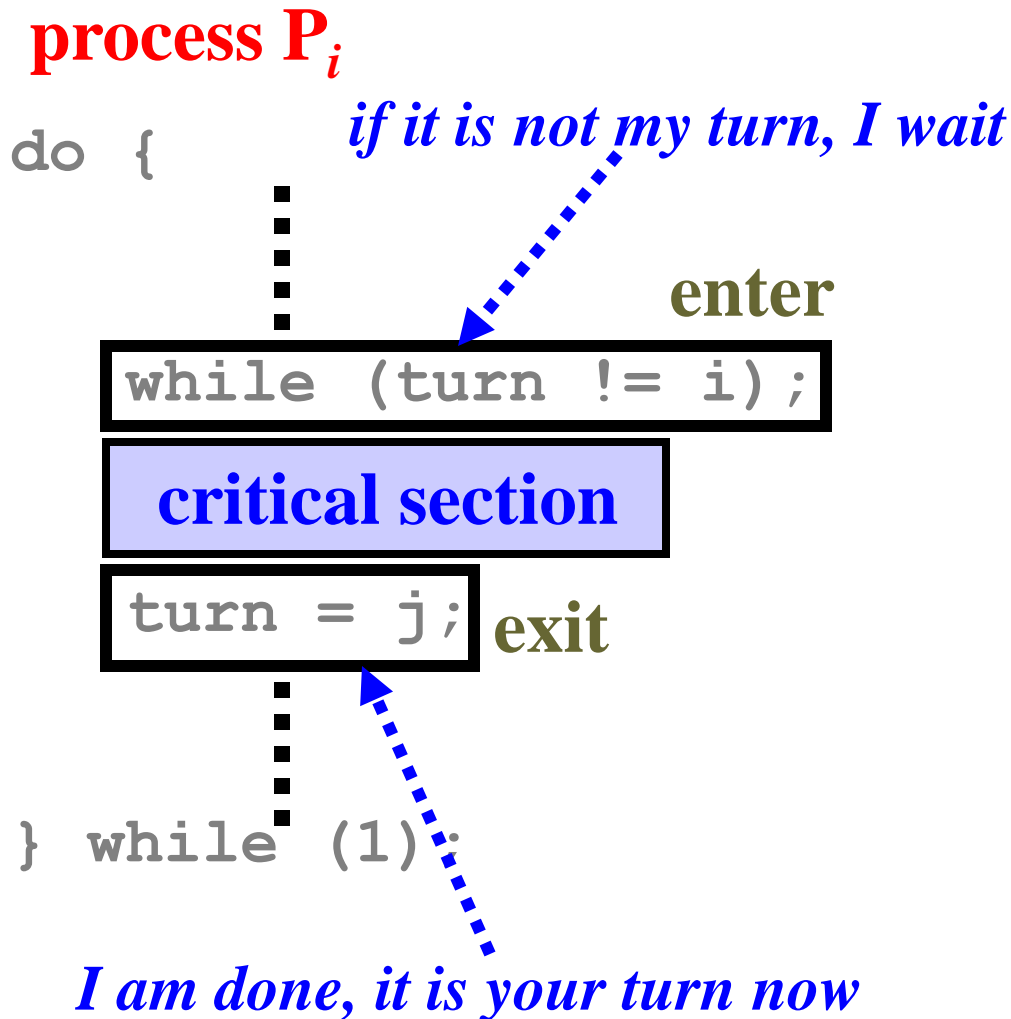
1

Pablo Picasso

Software Solutions for Two Processes

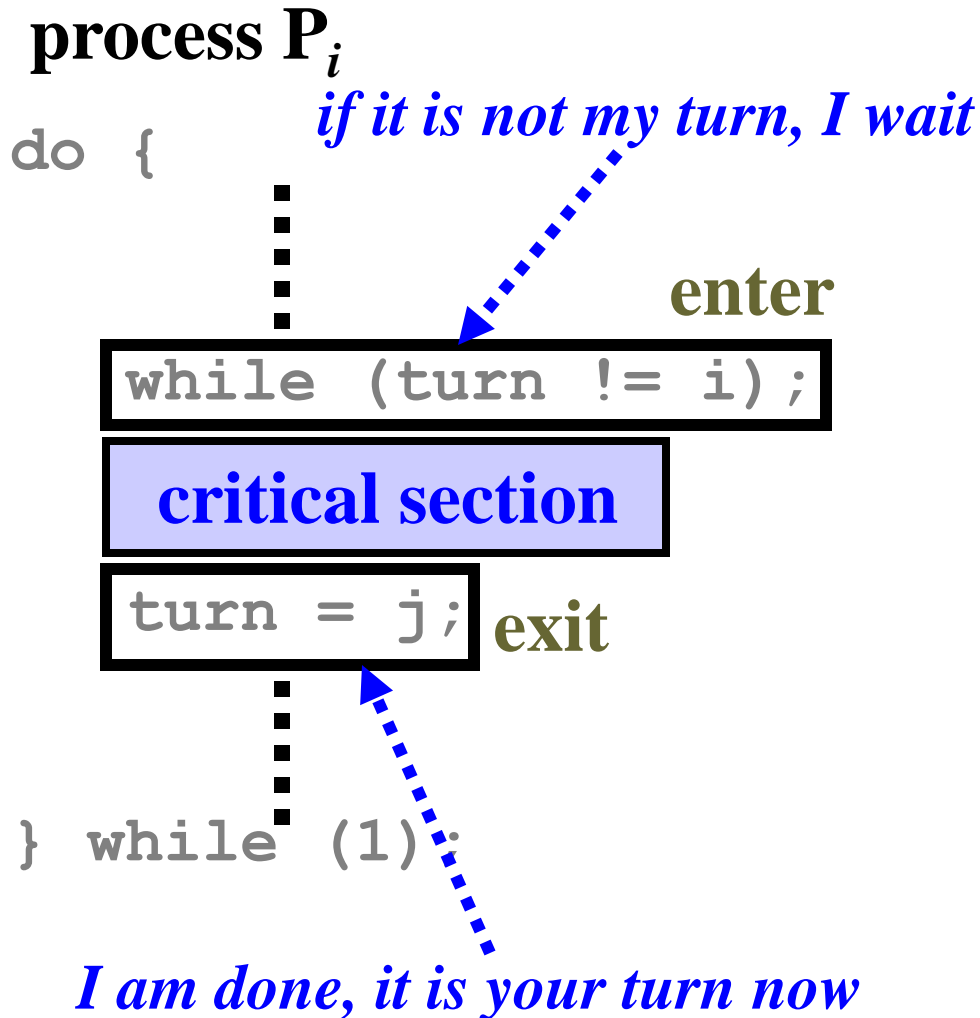
- Suppose we have two processes P_0 and P_1 .
- Let one process be P_i and the other be P_j , where $j = 1 - i$. Thus, if $i = 0$, then $j = 1$ and if $i = 1$, then $j = 0$.
- We wish to design an enter-exit protocol for a critical section to ensure mutual exclusion.
- We will go through a few unsuccessful attempts and finally obtain a correct one.
- These solutions are pure software-based.

Attempt I: 1/3



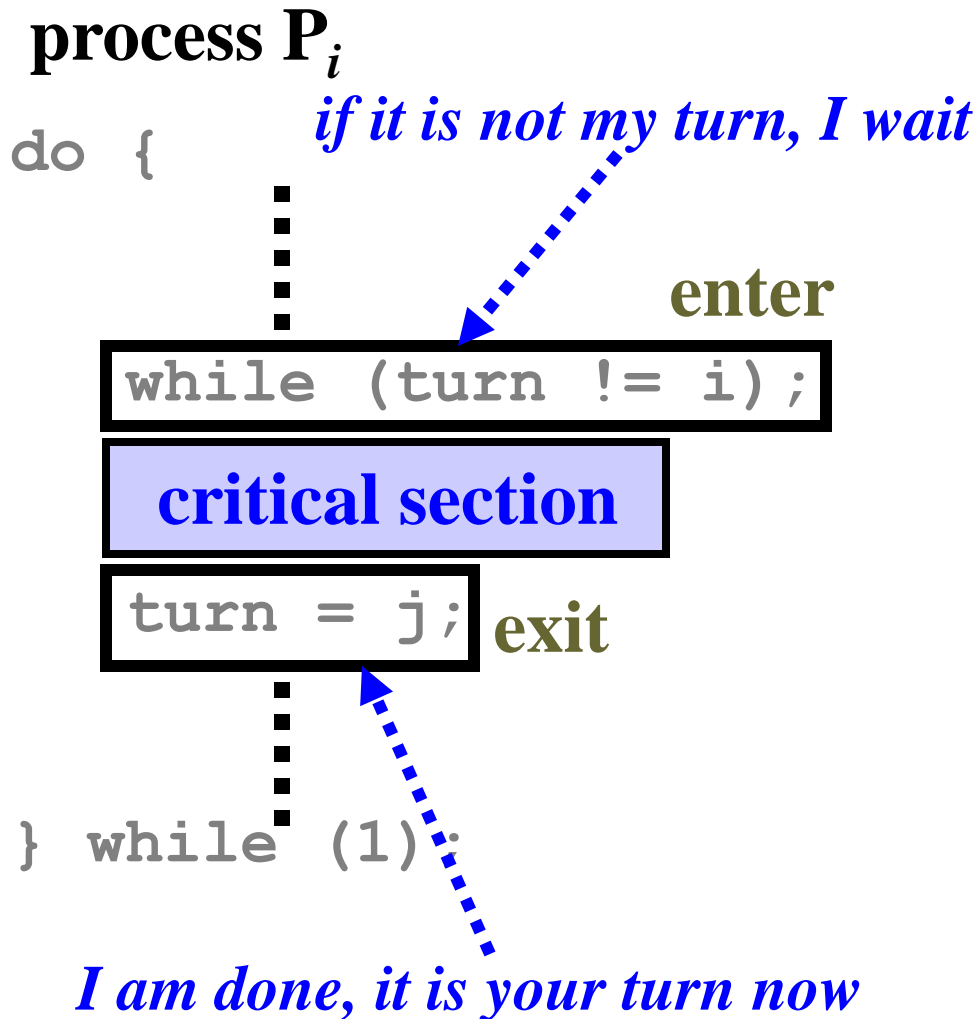
- Shared variable `turn` controls who can enter the critical section.
- Since `turn` is either 0 or 1, only one can enter.
- However, processes are forced to run in an *alternating* way.
- **Not good!**

Attempt I: 2/3



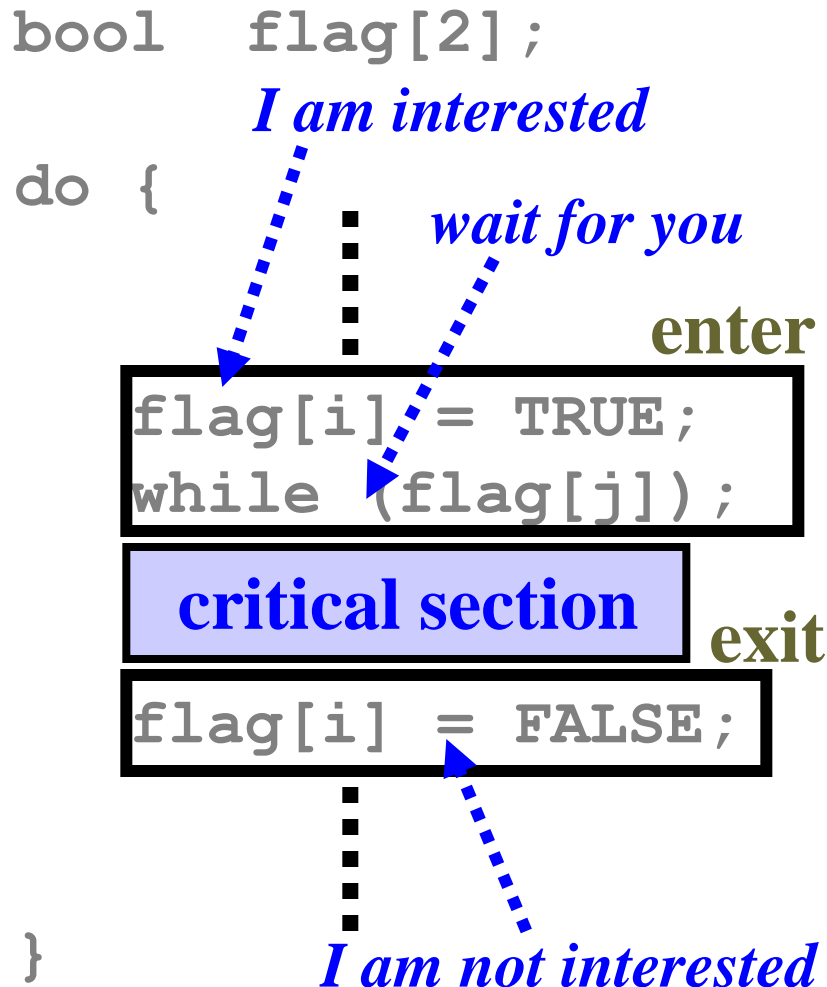
- **Mutual Exclusion**
- P_0 in its CS if $\text{turn}=0$.
- P_1 in its CS if $\text{turn}=1$.
- If P_0 and P_1 are **BOTH** in their CS, then $\text{turn}=0$ and $\text{turn}=1$ must **BOTH** be true.
- This is absurd, because a variable can only hold one and only one value (i.e., cannot hold both 0 and 1).

Attempt I: 3/3



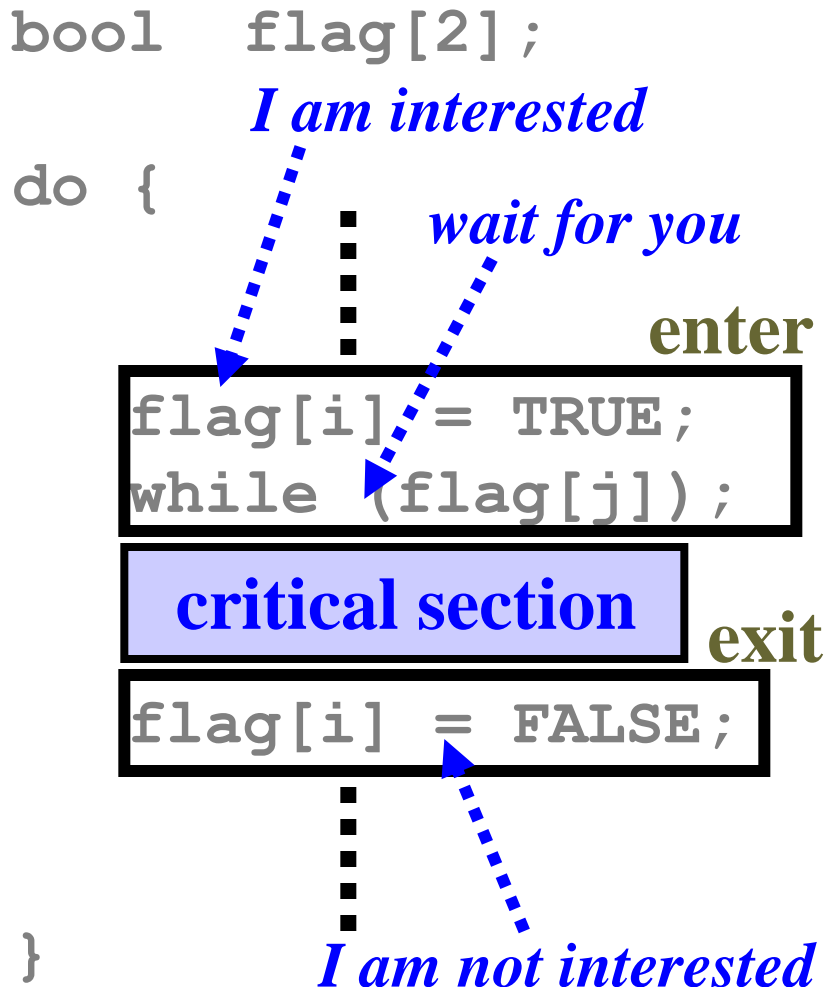
- **Progress**
- If P_i sets `turn` to j and never uses the critical section again, P_j can enter but cannot enter again.
- Thus, an irrelevant process blocks other processes from entering a critical section. **Not good!**
- Does bounded waiting hold? **Exercise!**⁵

Attempt II: 1/5



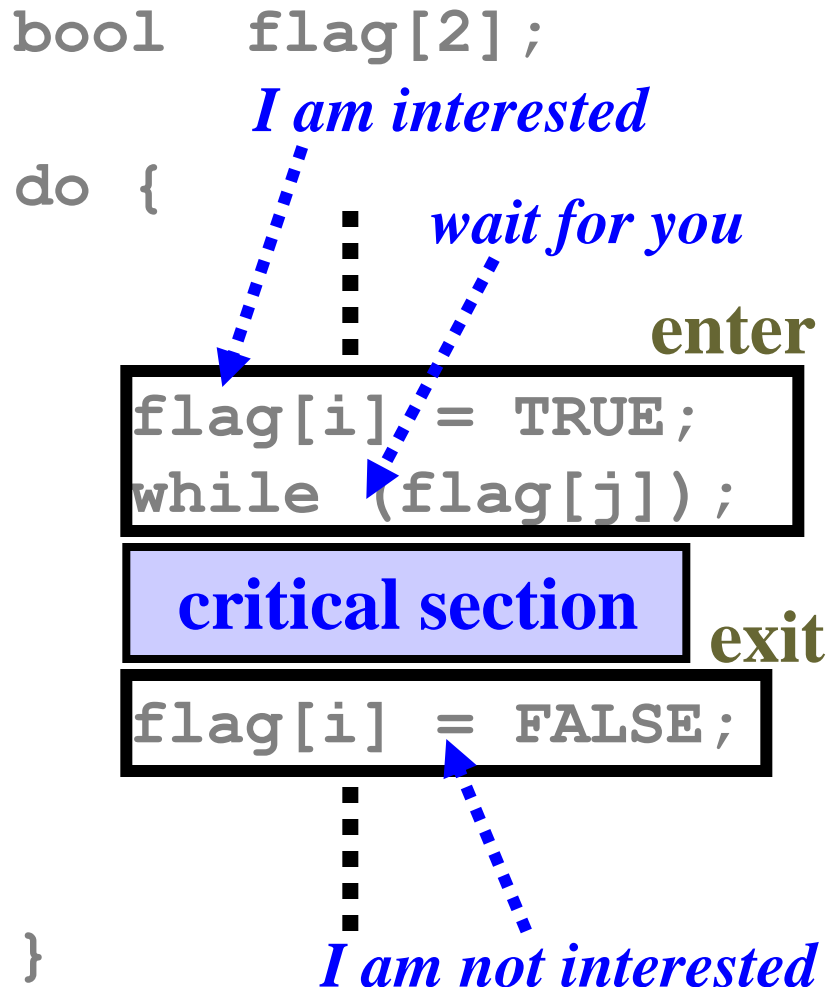
- Shared variable `flag[i]` is the “state” of process P_i : *interested* or *not-interested*.
- P_i indicates its intention to enter, waits for P_j to exit, enters its section, and, finally, changes to “*I am out*” upon exit.

Attempt II: 2/5



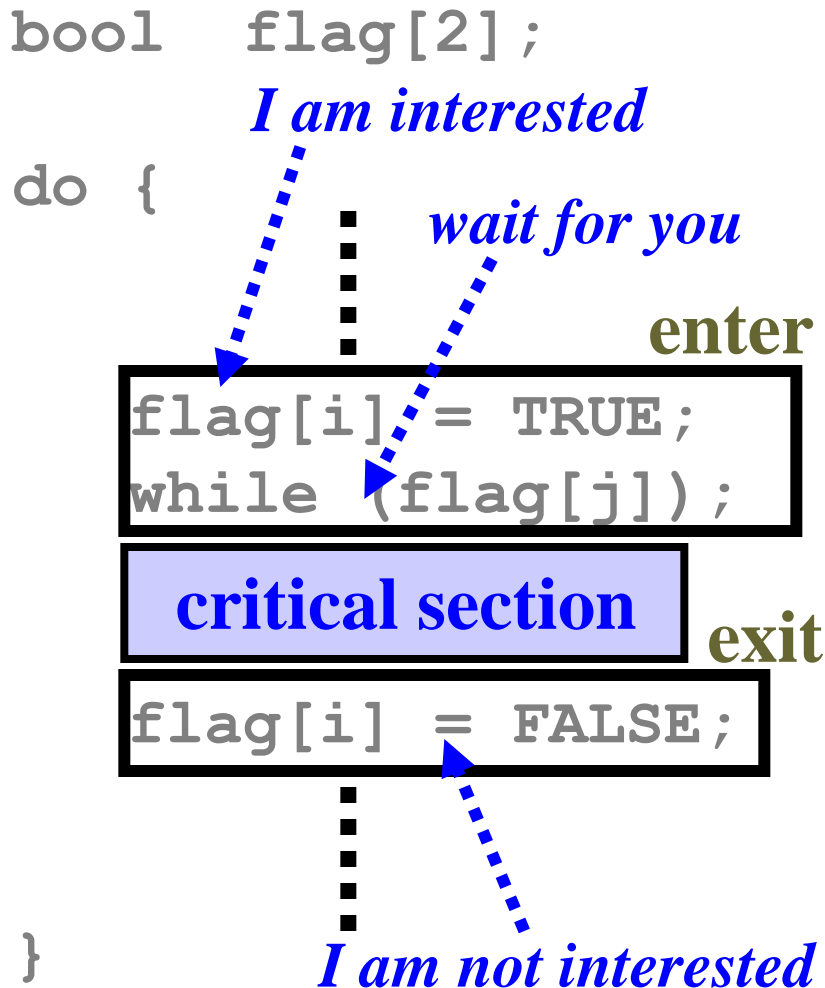
- **Mutual Exclusion**
- P_0 is in CS if `flag[0]` is TRUE **AND** `flag[1]` is FALSE.
- P_1 is in CS if `flag[1]` is TRUE **AND** `flag[0]` is FALSE.
- If both are in their CS, `flag[0]` must be both TRUE and FALSE.
- This is absurd.

Attempt II: 3/5



- **Progress**
- If both P_0 and P_1 set `flag[0]` and `flag[1]` to TRUE at the same time, then both will loop at the `while` forever and no one can enter.
- A decision cannot be made in finite time.

Attempt II: 4/5



- **Bounded Waiting**
- P_0 is in the enter section but switched out before setting `flag[0]` to TRUE.
- P_1 reaches its CS and sees `flag[0]` being not TRUE. P_1 enters.
- P_1 can enter and exit in this way repeatedly. Thus, P_0 cannot enter forever (i.e., unbounded).

Attempt II: 5/5

```

bool flag[2];

do {
    I am interested
    ...
    wait for you
    enter
    flag[i] = TRUE;
    while (flag[j]);
    critical section
    exit
    flag[i] = FALSE;
    ...
    I am not interested
}

```

no change to continue

P ₀	P ₁	flag[0]	flag[1]
LOAD #0		F	F
	flag[1]=T	F	T
	while ..	F	T
	in CS		
	flag[i]=F	F	F
	loop back		

```

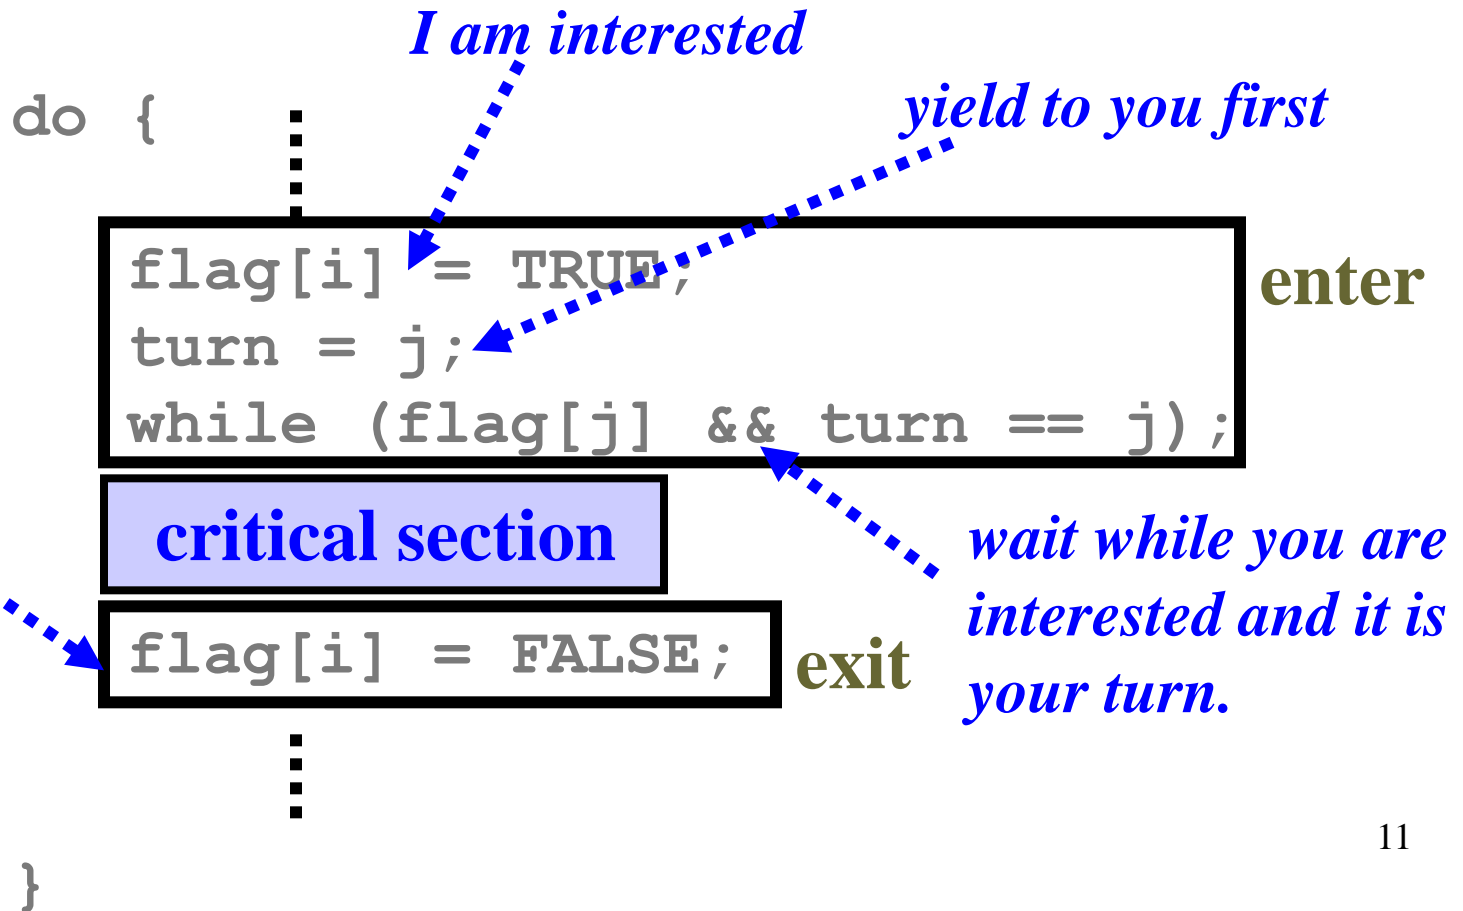
flag[i] = LOAD i
          LOAD address flag[i]
          MOVE T or F to flag[i]

```

Attempt III/A Combination: 1/12

Peterson's Algorithm

```
bool flag[2] = FALSE; // process  $P_i$   
int  turn;
```



Attempt III: Mutual Exclusion 2/12

process P_i

```
flag[i] = TRUE;  
turn = j;  
while (flag[j] && turn == j);
```

process P_j

```
flag[j] = TRUE;  
turn = i;  
while (flag[i] && turn == i);
```

- If P_i is in its critical section, then it sets
 - ❖ `flag[i]` to TRUE
 - ❖ `turn` to j (but `turn` may not be j after this point because P_j **may** set it to i later).
 - ❖ and waits until `flag[j] && turn == j` becomes FALSE

Attempt III: Mutual Exclusion 3/12

process P_i

```
flag[i] = TRUE;  
turn = j;  
while (flag[j] && turn == j);
```

process P_j

```
flag[j] = TRUE;  
turn = i;  
while (flag[i] && turn == i);
```

- If P_j is in its critical section, then it sets
 - ❖ `flag[j]` to TRUE
 - ❖ `turn` to `i` (but `turn` may not be `i` after this point because P_i **may** set it to `j` later).
 - ❖ and waits until `flag[i] && turn == i` becomes FALSE

Attempt III: Mutual Exclusion 4/12

process P_i

```
flag[i] = TRUE;  
turn = j;  
while (flag[j] && turn == j);
```

process P_j

```
flag[j] = TRUE;  
turn = i;  
while (flag[i] && turn == i);
```

- If processes P_i and P_j are both in their critical sections, then we have:

❖ $flag[i]$ and $flag[j]$ are both TRUE. they are both TRUE

❖ $flag[i] \&\& turn == i$ and $flag[j] \&\& turn == j$ are both FALSE.

- ❖ Therefore, $turn == i$ and $turn == j$ must both be FALSE.

Attempt III: Mutual Exclusion 5/12

process P_i

```
flag[i] = TRUE;  
turn = j;  
while (flag[j] && turn == j);
```

process P_j

```
flag[j] = TRUE;  
turn = i;  
while (flag[i] && turn == i);
```

- Since $\text{turn} == i$ and $\text{turn} == j$ are both FALSE and since turn is set to j (by P_i) or i (by P_j) before entering the critical section, only one of $\text{turn} == i$ and $\text{turn} == j$ can be FALSE but not both.
- Therefore, we have a contradiction.
- Of course, mutual exclusion holds.

Attempt III: Mutual Exclusion 6/12

- We normally use the proof by contradiction technique to establish the mutual exclusion condition.
- To do so, follow the procedure below:
 - ❖ Find the condition C_0 for P_0 to enter its CS
 - ❖ Find the condition C_1 for P_1 to enter its CS
 - ❖ If P_0 and P_1 are in their critical sections, C_0 and C_1 will both be true.
 - ❖ From C_0 and C_1 being true, we should be able to derive an absurd result.
 - ❖ Therefore, mutual exclusion holds.

Attempt III: Mutual Exclusion 7/12

- We care about the conditions C_0 and C_1 . The way of reaching these conditions via instruction execution is un-important.
- Never use an execution sequence to prove mutual exclusion. In doing so, you make a serious mistake, which is usually referred to as **proof by example**.
- You may use a single example to show a proposition being false. But, you cannot use a single example to show a proposition being true. That is, $3^2 + 4^2 = 5^2$ cannot be used to prove $a^2 + b^2 = c^2$ for a right triangles.

Attempt III: Progress 8/12

process P_i

```
flag[i] = TRUE;  
turn = j;  
while (flag[j] && turn == j);
```

process P_j

```
flag[j] = TRUE;  
turn = i;  
while (flag[i] && turn == i);
```

- If P_i and P_j are both waiting to enter their critical sections, since the value of `turn` can only be i or j but not both, one process can pass its `while` loop (*i.e.*, decision time is finite).
- Suppose P_i is waiting and P_j is not in its CS:
 - ❖ Since P_j is **not interested** in entering, `flag[j]` was set to `FALSE` when P_j exits and P_i enters.
 - ❖ Thus, the process that is not entering does not influence the decision.

Attempt III: Bounded Waiting 9/12

process P_i

```
flag[i] = TRUE;  
turn = j;  
while (flag[j] && turn == j);
```

process P_j

```
flag[j] = TRUE;  
turn = i;  
while (flag[i] && turn == i);
```

- If P_i wishes to enter, we have three cases:
 1. P_j is *outside* of its critical section.
 2. P_j is *in* its critical section.
 3. P_j is *in the entry section*.

Attempt III: Bounded Waiting 10/12

process P_i

```
flag[i] = TRUE;  
turn = j;  
while (flag[j] && turn == j);
```

process P_j

```
flag[j] = TRUE;  
turn = i;  
while (flag[i] && turn == i);
```

- **CASE I**: If P_j is *outside* of its critical section, P_j sets `flag[j]` to FALSE when it exits its critical section, and P_i may enter.
- In this case, P_i does not wait.

Attempt III: Bounded Waiting 11/12

process P_i

```
flag[i] = TRUE;  
turn = j;  
while (flag[j] && turn == j);
```

process P_j

```
flag[j] = TRUE;  
turn = i;  
while (flag[i] && turn == i);
```

- **CASE 2**: If P_j is *in the entry section*, depending on the value of `turn`, we have two cases:
 - ❖ If `turn` is `i` (e.g., P_i sets `turn` to `j` before P_j sets `turn` to `i`), P_i enters immediately.
 - ❖ Otherwise, P_j enters and P_i stays in the `while` loop, and we have **CASE 3**.

Attempt III: Bounded Waiting 12/12

process P_i

process P_j

<pre>flag[i] = TRUE; turn = j; while (flag[j] && turn == j);</pre>	<pre>flag[j] = TRUE; turn = i; while (flag[i] && turn == i);</pre>
--	--

- **CASE 3:** If P_j is *in* its critical section, turn must be j and P_i waits for at most one round.

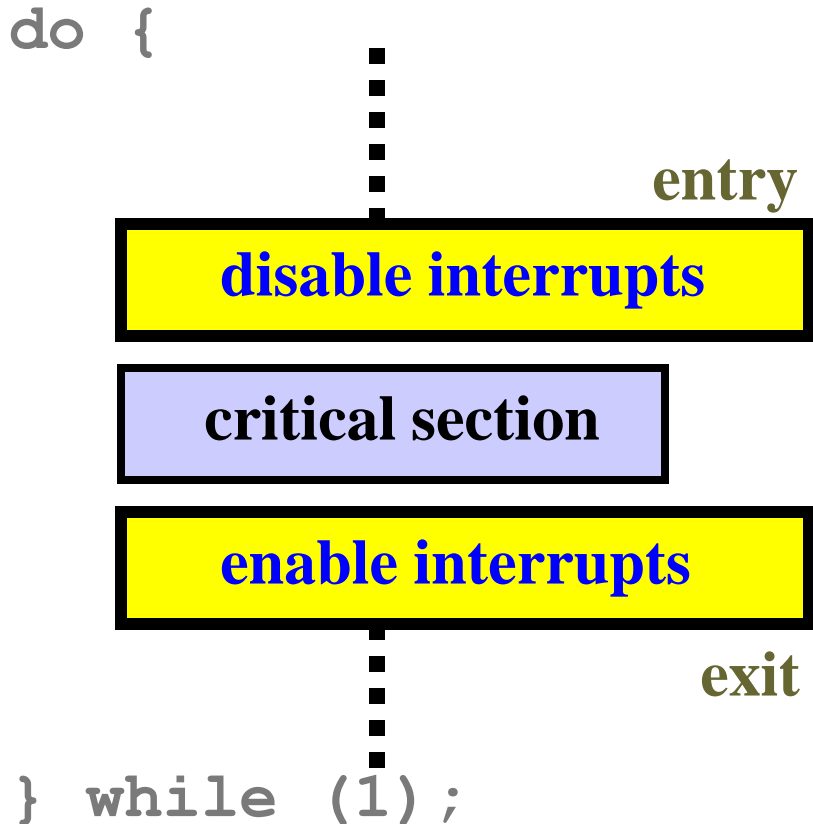
P_i	P_j	flag[i]	flag[j]	turn	Comments
flag[i]=T	flag[j]=T	TRUE	TRUE	?	
while (...)		TRUE	TRUE	j	P_j enters
	Critical Sec				P_j in CS
	flag[j]=F	TRUE	FALSE	j	P_j exits
	flag[j]=T	TRUE	TRUE	j	P_j returns
	turn = i	TRUE	TRUE	i	P_j yields
	while (...)	TRUE	TRUE	i	P_j loops
Critical Sec					P_i enters

P_i has a chance to enter here.
if P_j comes back fast

Hardware Support

- There are two types of hardware synchronization supports:
 - ❖ Disabling/Enabling interrupts: This is slow and difficult to implement on multiprocessor systems.
 - ❖ Special *privileged*, actually *atomic*, machine instructions:
 - ✓ Test and set (TS)
 - ✓ Swap
 - ✓ Compare and Swap (CS)

Interrupt Disabling



- Because interrupts are disabled, no **context switch** can occur in a critical section (why?).
- Infeasible in a **multiprocessor** system because all CPUs/cores must be informed.
- Some features that depend on interrupts (*e.g.*, clock) may not work properly.

Test-and-Set Instruction: 1/2

```
bool TS(bool *key)
{
    bool save = *key;
    *key = TRUE;
    return save;
}
```

- **TS** is atomic.
- **Mutual exclusion** is met as the TS instruction is atomic. **See next slide.**
- However, **bounded waiting** may not be satisfied. **Progress?**

```
bool lock = FALSE;

do {
    ⋮
    entry
    while (TS(&lock));
    critical section
    lock = FALSE; exit
    ⋮
} while (1);
```

Test-and-Set Instruction: 2/2

- A process is in its critical section if the TS instruction returns **FALSE**.
- If two processes P_0 and P_1 are in their critical sections, they both got the **FALSE** return value from TS.
- P_0 and P_1 cannot execute their TS instructions at the same time because TS is atomic.
- Hence, one of them, say P_0 , executes the TS instruction before the other.
- Once P_0 finishes its TS, the value of `lock` becomes **TRUE**.
- P_1 cannot get a **FALSE** return value and cannot enter its CS.
- We have a contradiction!

```
bool lock = FALSE;

do {
    ■
    ■
    ■
    while (TS(&lock));
    critical section
    lock = FALSE;
    ■
    ■
    ■
} while (1);
```

Problems with Software and Hardware Solutions

- All of these solutions use **busy waiting**.
- **Busy waiting** means a process waits by executing a tight loop to check the status/value of a variable.
- Busy waiting may be needed on a multiprocessor system; however, it wastes CPU cycles that some other processes may use productively.
- Even though some systems may allow users to use some atomic instructions, unless the system is lightly loaded, CPU and system performance can be low, although a programmer may “think” his/her program looks more efficient.
- So, we need better solutions.

The End