

# Part III

# Synchronization

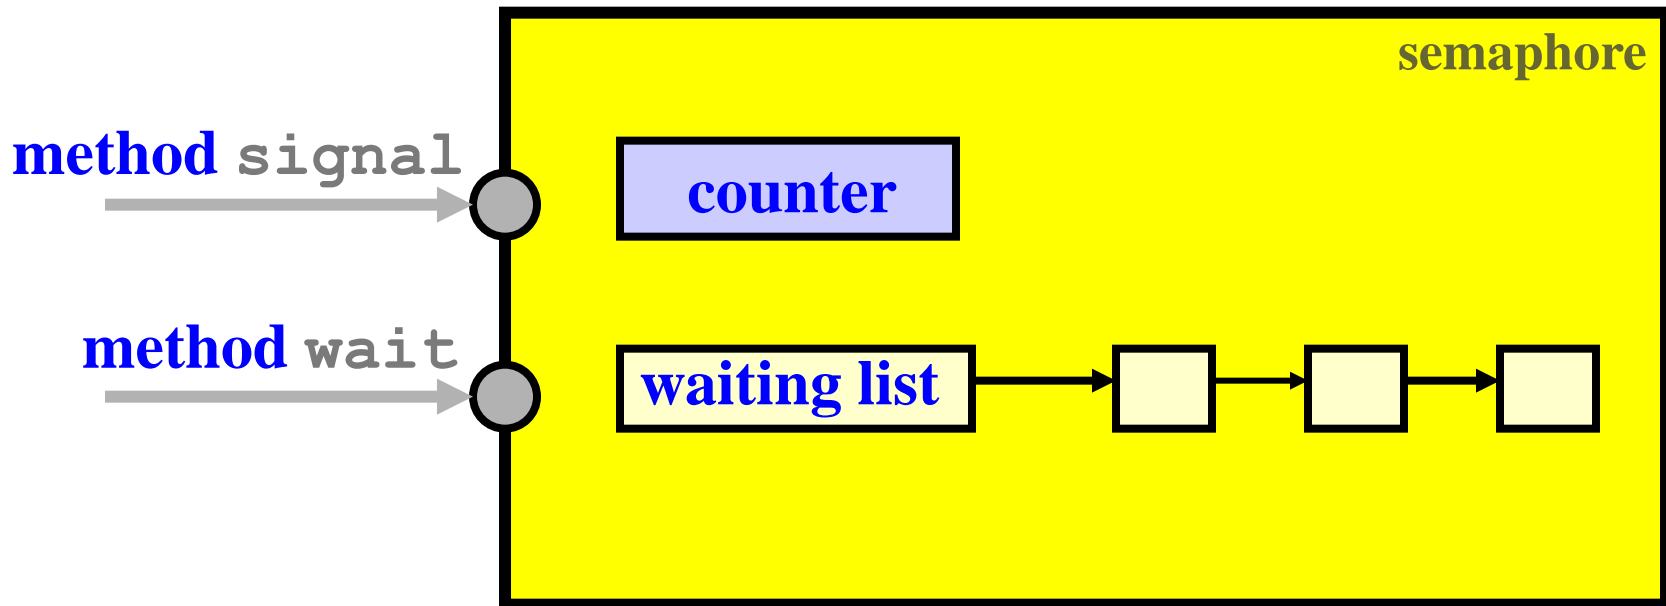
## Semaphores

*The bearing of a child takes nine months,  
no matter how many women are assigned.*

1

# Semaphores

- A *semaphore* is an object that consists of a private **counter**, a private **waiting list** of processes, and two public **methods** (e.g., member functions): **signal** and **wait**.



# Semaphore Method: wait

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting list;
        block();
    }
}
```

- After decreasing the counter by 1, if the new value becomes negative, then
  - ❖ add the caller to the waiting list, and
  - ❖ block the caller.

# Semaphore Method: signal

```
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a process P from the waiting list;
        resume (P) ;
    }
}
```

- After increasing the counter by 1, if the new value is not positive (e.g., non-negative), then
  - ❖ remove a process **P** from the waiting list,
  - ❖ resume the execution of process **P**, and return

# Important Note: 1/4

<pre>S.count--; if (S.count&lt;0) {     add to list;     block(); }</pre>	<pre>S.count++; if (S.count&lt;=0) {     remove P;     resume(P); }</pre>
---	---

- If  $S.count < 0$ ,  $abs(S.count)$  is the number of waiting processes.
- This is because processes are added to (*resp.*, removed from) the waiting list only if the counter value is  $< 0$  (*resp.*,  $\leq 0$ ).

# Important Note: 2/4

<pre>S.count--; if (S.count&lt;0) {     add to list;     block(); }</pre>	<pre>S.count++; if (S.count&lt;=0) {     remove P;     resume(P); }</pre>
---	---

- The waiting list can be implemented with a queue if FIFO order is desired.
- However, **the correctness of a program should not depend on a particular implementation (e.g., ordering) of the waiting list.**

# Important Note: 3/4

<pre>S.count--; if (S.count&lt;0) {     add to list;     block(); }</pre>	<pre>S.count++; if (S.count&lt;=0) {     remove P;     resume(P); }</pre>
---	---

- The caller may block in the call to `wait()`.
- The caller never blocks in the call to `signal()`.  
If `S.count > 0`, `signal()` returns and the caller continues. Otherwise, a waiting process is released and the caller continues. In this case, **two** processes continue.

# The Most Important Note: 4/4

<pre>S.count--; if (S.count&lt;0) {     add to list;     block(); }</pre>	<pre>S.count++; if (S.count&lt;=0) {     remove P;     resume(P); }</pre>
---	---

- `wait()` and `signal()` must be executed **atomically** (*i.e.*, as one **uninterruptible** unit).
- Otherwise, **race conditions** may occur.
- **Homework:** use execution sequences to show race conditions if `wait()` and/or `signal()` is not executed atomically.



# Typical Uses of Semaphores

- There are three typical uses of semaphores:

- ❖ **mutual exclusion:**

- Mutex (*i.e.*, *Mutual Exclusion*) locks

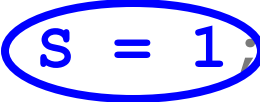
- ❖ **count-down lock:**

- Keep in mind that a semaphore has a private counter that can count.

- ❖ **notification:**

- Wait for an event to occur and indicate an event has occurred.

# Use 1: Mutual Exclusion (Lock)

semaphore **S = 1;**  initialization is important

```
int      count = 0;    // shared variable
```

Process 1	Process 2
while (1) {	while (1) {
// do something entry	// do something
<b>S.wait();</b>	<b>S.wait();</b>
<b>count++;</b> critical sections	<b>count--;</b>
<b>S.signal();</b>	<b>S.signal();</b>
// do something exit	// do something
}	}

- What if the initial value of S is zero?
- S is a *binary semaphore* (count being 0 or 1).

## Use 2: Count-Down Counter

semaphore S = 3;

Process 1

```
while (1) {  
    // do something  
    S.wait();  
    S.signal();  
    // do something  
}
```

Process 2

```
while (1) {  
    // do something  
    S.wait();  
    S.signal();  
    // do something  
}
```

at most 3 processes can be here!!!

- After three processes pass through `wait()`, this section is locked until a process calls `signal()`.

# Use 3: Notification


```
semaphore S1 = 1, S2 = 0;
```

process 1		process 2
<pre>while (1) {     // do something     S1.wait();     cout &lt;&lt; "1";     S2.signal();     // do something }</pre>	<pre>graph LR; subgraph P1 [process 1]; S2s[S2.signal()]; end; subgraph P2 [process 2]; S1w[S1.wait()]; S1s[S1.signal()]; end; S2s -.-&gt; notify  S1w; S1s -.-&gt; notify  S2w[S2.wait()];</pre>	<pre>while (1) {     // do something     S2.wait();     cout &lt;&lt; "2";     S1.signal();     // do something }</pre>

- Process 1 uses `S2.signal()` to notify process 2, indicating **“I am done. Please go ahead.”**
- The output is 1 2 1 2 1 2 .....
- What if `S1` and `S2` are both 0's or both 1's?
- What if `S1 = 0` and `S2 = 1`?

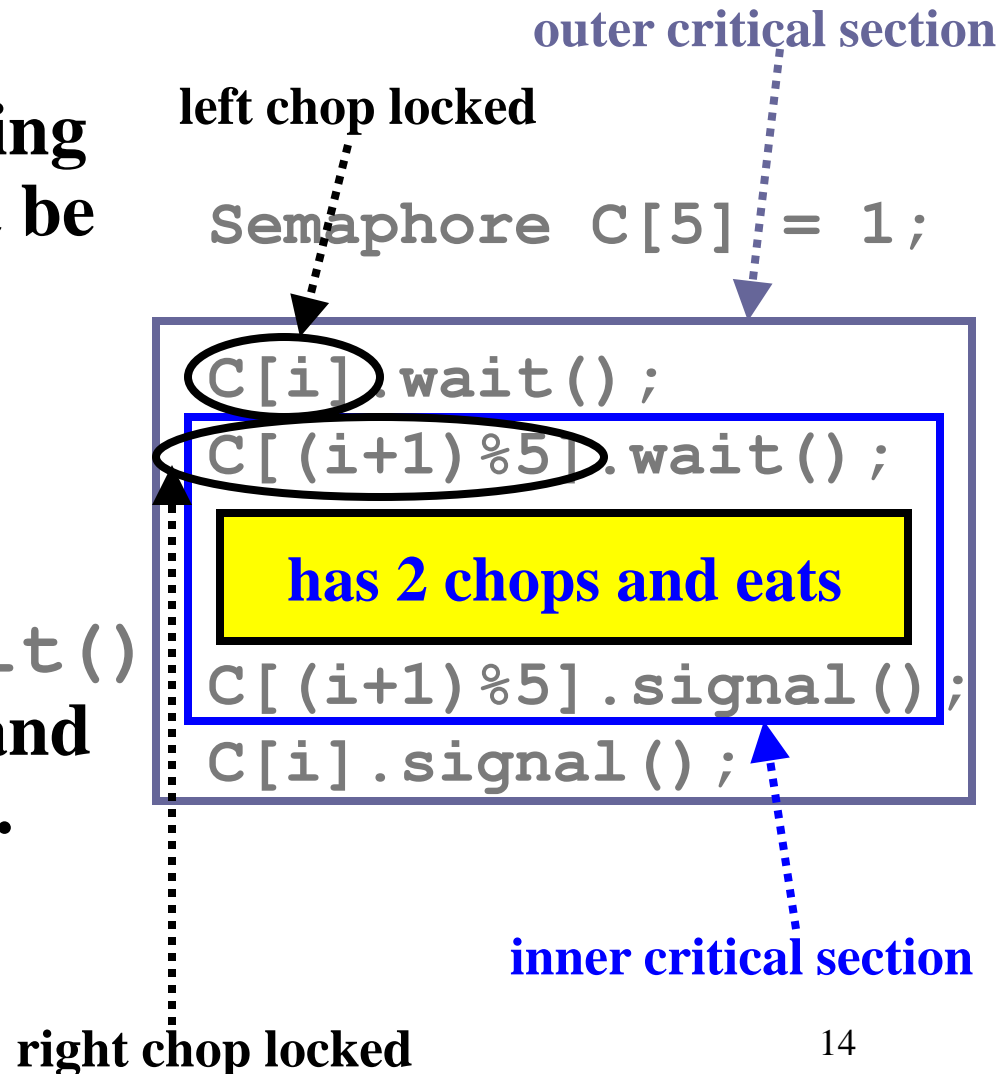
# Dining Philosophers

- Five philosophers are in a **thinking** - **eating** cycle.
- When a philosopher gets hungry, he sits down, picks up *his left* and then *his right* chopsticks, and eats.
- A philosopher can eat only if he has **both** chopsticks.
- After eating, he puts down both chopsticks and thinks.
- This cycle continues.



# Dining Philosopher: Ideas

- Chopsticks are shared items (by two neighboring philosophers) and must be protected.
- Each chopstick has a semaphore with initial value 1 (i.e., available).
- A philosopher calls `wait()` to pick up a chopstick and `signal()` to release it.



# Dining Philosophers: Code

```
semaphore C[5] = 1;
```

**philosopher  $i$**

```
while (1) {  
    // thinking  
    C[i].wait();  
    C[(i+1)%5].wait();  
    // eating  
    C[(i+1)%5].signal();  
    C[i].signal();  
    // finishes eating  
}
```

wait for my left chop

wait for my right chop


release my right chop

release my left chop

**Does this solution work?**

# Dining Philosophers: Deadlock!

- If all five philosophers sit down and pick up their left chopsticks at the same time, this causes a **circular waiting** and the program deadlocks.
- An easy way to remove this deadlock is to introduce a weirdo who picks up his **right** chopstick first!





# Dining Philosophers: A Better Idea

```
semaphore C[5] = 1;
```

philosopher  $i$  (0, 1, 2, 3)

```
while (1) {  
    // thinking  
    C[i].wait();  
    C[(i+1)%5].wait();  
    // eating  
    C[(i+1)%5].signal();  
    C[i].signal();  
    // finishes eating;  
}
```

lock left chop

Philosopher 4: the weirdo


```
while (1) {  
    // thinking  
    C[(i+1)%5].wait();  
    C[i].wait();  
    // eating  
    C[i].signal();  
    C[(i+1)%5].signal();  
    // finishes eating  
}
```

lock right chop

# Dining Philosophers: Questions

- The following are some important questions for you to think about.
  - ❖ We choose philosopher 4 to be the weirdo. Does this choice matter?
  - ❖ Show that this solution does not cause **circular waiting**.
  - ❖ Show that this solution does not cause **circular waiting** even if we have more than 1 and less than 5 weirdoes.
- This solution is not **symmetric** because not all threads run the same code.

# Count-Down Lock Example



- The naïve solution to the dining philosophers problem causes circular waiting.
- If only **four** philosophers are allowed to sit down, deadlock cannot occur.
- **Why?** If all four of them sit down at the same time, the right-most philosopher may have both chopsticks!
- **How about fewer than four?** This is obvious.

# Count-Down Lock Example

```
semaphore C[5]= 1;
```

```
semaphore Chair = 4;
```

get a chair

```
while (1) {
```

```
    // thinking
```

```
    Chair.wait();
```

```
        C[i].wait();
```

```
        C[(i+1)%5].wait();
```

```
        // eating
```

```
        C[(i+1)%5].signal();
```

```
        C[i].signal();
```

```
    Chair.signal();
```


```
}
```

this is a count-down lock  
that only allows 4 to go!

this is our old friend


release my chair

# The Producer/Consumer Problem




- Suppose we have a **circular buffer** of  $n$  slots.
- Pointer ***in*** (*resp.*, ***out***) points to the first **empty** (*resp.*, **filled**) slot.
- **Producer** processes keep adding data into the buffer.
- **Consumer** processes keep retrieving data from the buffer.

# Problem Analysis



*buffer is implemented  
with an array `Buf [ ]`*

- A producer deposits data into `Buf[in]` and a consumer retrieves info from `Buf[out]`.
- `in` and `out` must be advanced.
- `in` is shared among producers.
- `out` is shared among consumers.
- If `Buf` is full, producers should be blocked.
- If `Buf` is empty, consumers should be blocked.



- A semaphore to protect the buffer.
- The second semaphore to block producers if the buffer is full.
- The third semaphore to block consumers if the buffer is empty.

# Solution

number of slots

semaphore NotFull=n NotEmpty=0, Mutex=1;

producer

```
while (1) {  
    NotFull.wait();  
    Mutex.wait();  
    Buf[in] = x;  
    in = (in+1)%n;  
    Mutex.signal();  
    NotEmpty.signal();  
}
```

consumer

```
while (1) {  
    NotEmpty.wait();  
    Mutex.wait();  
    x = Buf[out];  
    out = (out+1)%n;  
    Mutex.signal();  
    NotFull.signal();  
}
```

notifications

critical section




# Question

- What if the producer code is modified as follows?
- **Answer:** a deadlock may occur. Why?

```
while (1) {  
    Mutex.wait();  
    NotFull.wait();  
    Buf[in] = x;  
    in = (in+1)%n;  
    NotEmpty.signal();  
    Mutex.signal();  
}
```

order changed

The diagram shows two blue dotted arrows originating from the text 'order changed'. One arrow points to the 'Mutex.wait()' line in the first box, and the other points to the 'Mutex.signal()' line in the second box, indicating a swap in their relative order in the original code.

# The Readers/Writers Problem

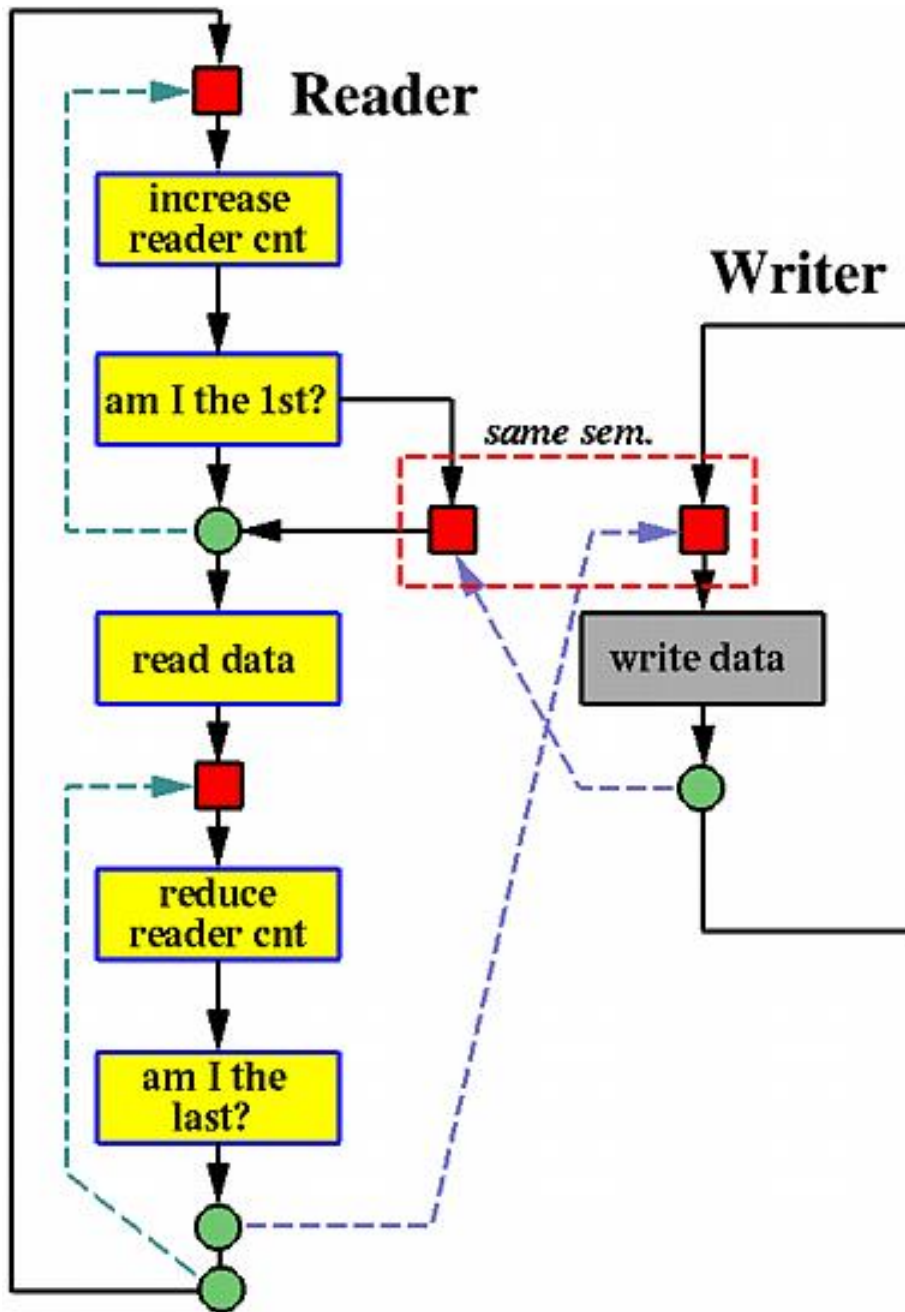
- Two groups of processes, **readers** and **writers**, access a shared resource by the following rules:
  - ❖ Readers can **read simultaneously**.
  - ❖ **Only one** writer can write at any time.
  - ❖ When a writer is writing, no reader can read.
  - ❖ If there is any reader reading, all **incoming writers must wait**. Thus, readers have higher priority.

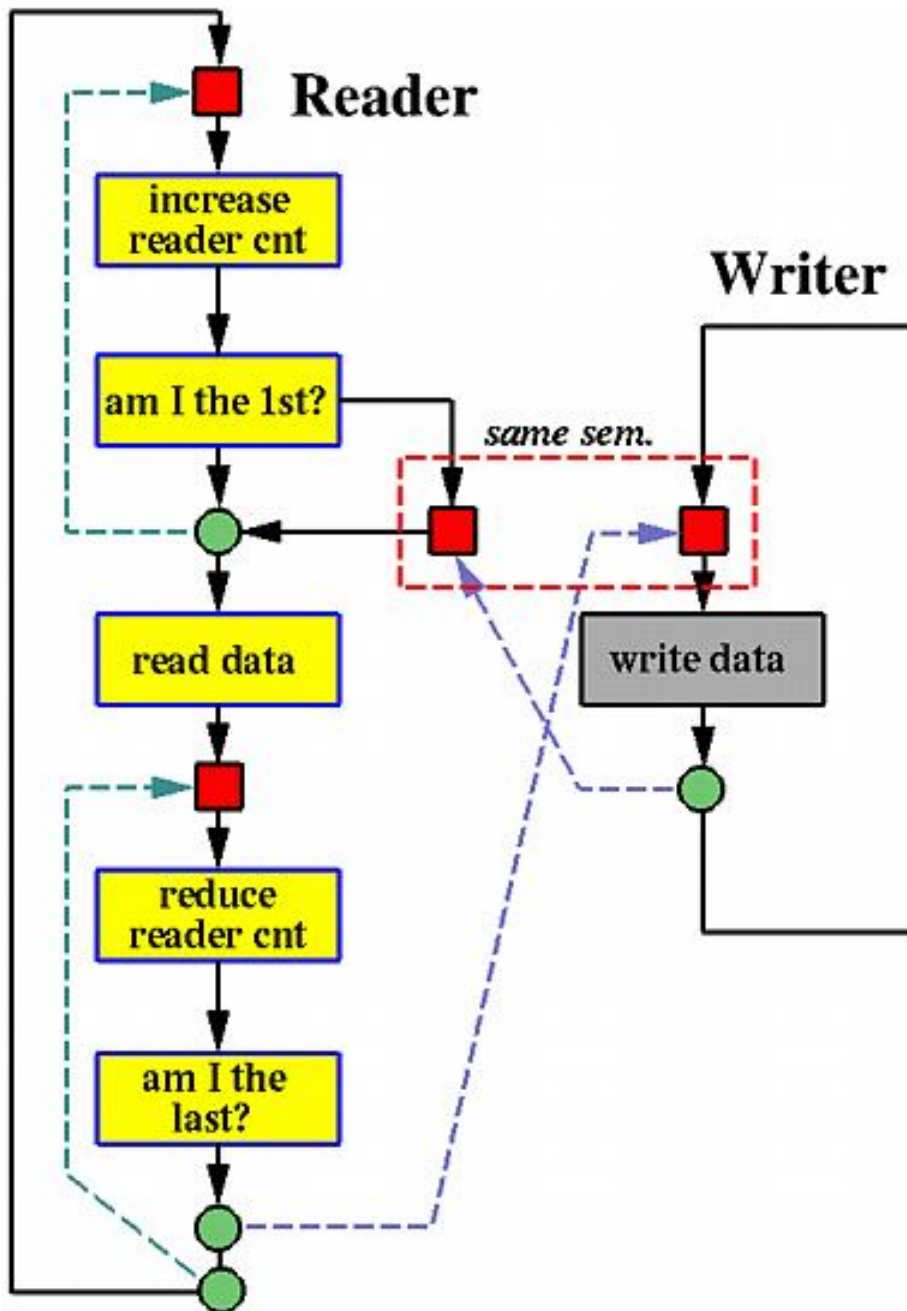
# Problem Analysis

- We need a semaphore to **block readers if a writer is writing**.
- When a writer arrives, it must **know if there are readers reading**. A reader count is required and must be protected by a lock.
- This **reader-priority** version has a problem: the bounded waiting condition may be violated if readers keep coming, causing the waiting writers no chance to write.

# Readers

- When a reader arrives, it adds 1 to the counter.
- If it is the first reader, waits until no writer is writing.
- Reads data.
- Decreases the counter.
- If it is the last reader, notifies the waiting readers and writers that no reader is reading.





# Writer

- When a writer comes in, it waits until no reader is reading and no writer is writing.
- Then, it writes data.
- Finally, notifies waiting readers and writers that no writer is writing.

# Solution

```
semaphore Mutex = 1, WrtMutex = 1;  
int          RdrCount;
```

## reader

```
while (1) {  
    Mutex.wait();  
    RdrCount++;  
    if (RdrCount == 1)  
        WrtMutex.wait();  
    Mutex.signal();  
    // read data  
    Mutex.wait();  
    RdrCount--;  
    if (RdrCount == 0)  
        WrtMutex.signal();  
    Mutex.signal();  
}
```

## writer

```
while (1) {  
    WrtMutex.wait();  
    // write data  
    WrtMutex.signal();  
}
```

**blocks both readers and writers**

# The Roller-Coaster Problem: 1/5

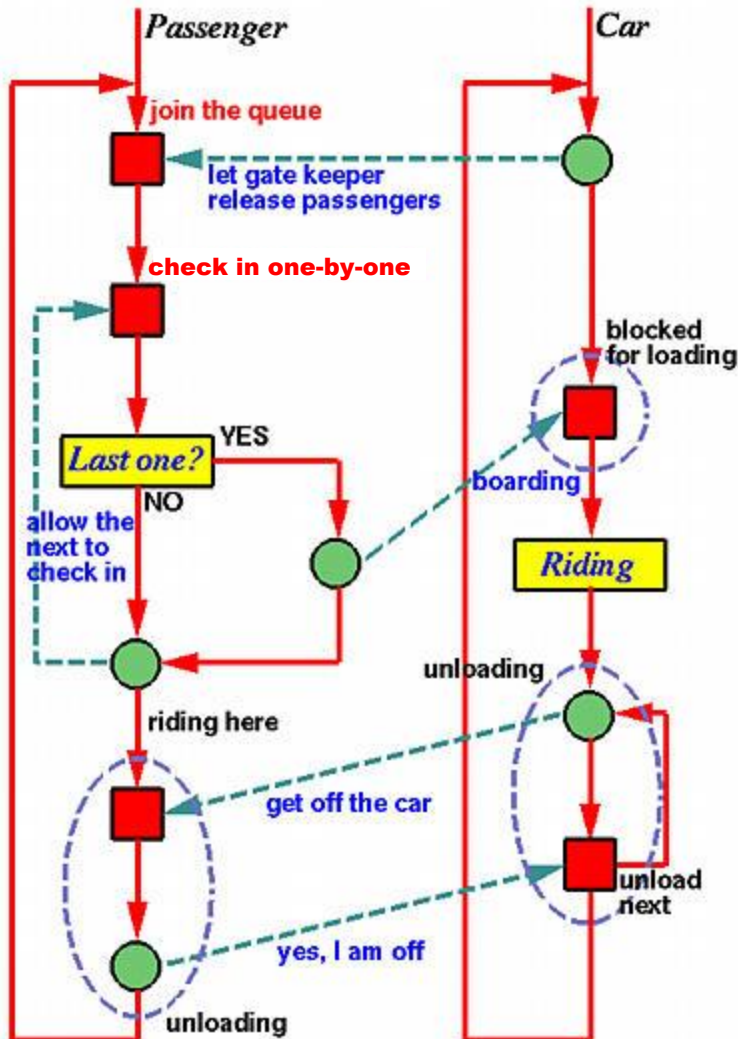
- Suppose there are  $n$  passengers and one roller coaster car. The passengers repeatedly wait to ride in the car, which can hold maximum  $C$  passengers, where  $C < n$ .
- The car can go around the track only when it is full. After finishes a ride, each passenger wanders around the amusement park before returning to the roller coaster for another ride.
- Due to safety reasons, the car only rides  $T$  times and then shut-down.

# The Roller-Coaster Problem: 2/5

- The car always rides with exactly  $C$  passengers
- No passengers will jump off the car while the car is running
- No passengers will jump on the car while the car is running
- No passengers will request another ride before they get off the car.

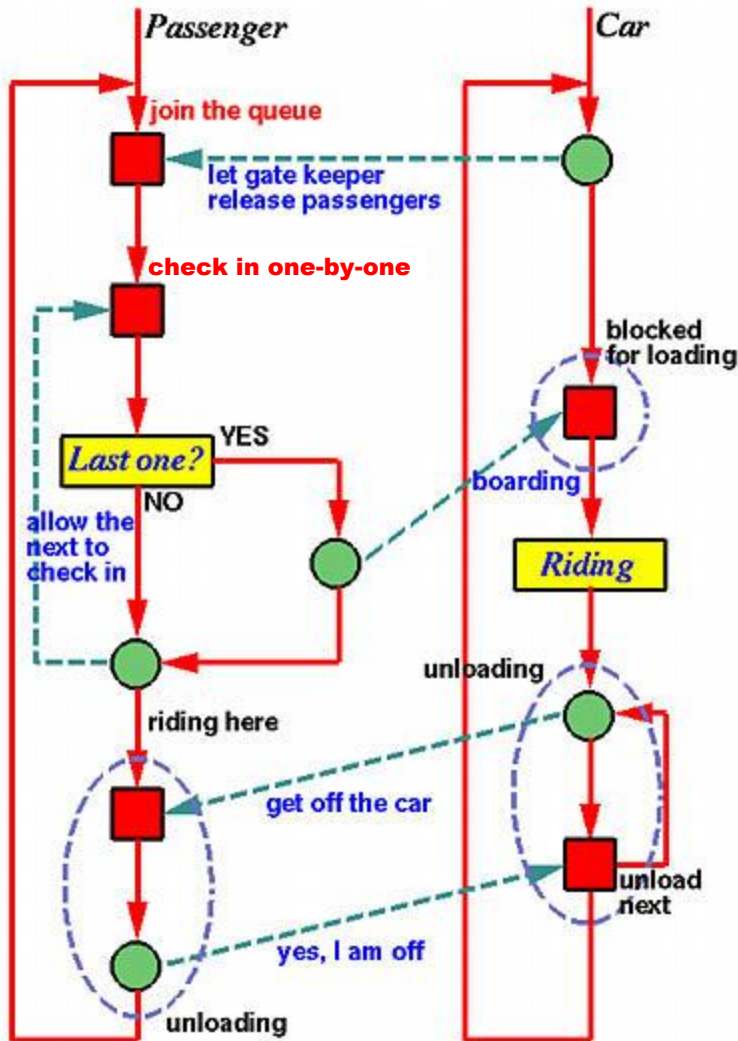


# The Roller-Coaster Problem: 3/5



- A **passenger** makes a decision to have a ride, and joins the queue.
- The queue is managed by a gate keeper.
- Passengers check in one-by-one.
- The last passenger tells the car that all passengers are on board.
- Then, they have a ride.
- After riding passengers get off the car one-by-one.
- They go back to play for a while and come back for a ride.

# The Roller-Coaster Problem: 4/5



- The car comes and lets the gate keeper know it is available so that the gate keeper could release passengers to check in.
- The car is blocked for loading.
- When the last passenger in the car, s/he informs the car that all passengers are on board, the car starts a ride.
- After this, the car waits until all passengers are off. Then, go for another round.

# The Roller-Coaster Problem: 5/5

```
int count = 0;  
Semaphore Queue = Boarding = Riding = Unloading = 0;  
Semaphore Check-In = 1;
```

## Passenger

```
Wait(Queue);  
Wait(Check-In);  
if (++count==Maximum)  
    Signal(Boarding);  
Signal(Check-In);  
Wait(Riding);  
Signal(Unloading);
```

## Car

```
for (i = 0; i < #rides; i++) {  
    count = 0; // reset counter before boarding  
    for (j = 1; j <= Maximum; j++)  
        Signal(Queue); // car available  
    Wait(Boarding);  
    // all passengers in car  
    // and riding  
    for (j = 1; j <= Maximum; j++) {  
        Signal(Riding);  
        Wait(Unloading);  
    }  
    // all passengers are off  
}
```

one ride

Unload passengers one-by-one  
Is this absolutely necessary?  
Can Unloading be removed? **Ex.**

# Semaphores with **ThreadMentor**

# Semaphores with ThreadMentor

- **ThreadMentor** has a class Semaphore with two methods `Wait()` and `Signal()`.
- Class Semaphore requires a non-negative integer as an initial value.
- A name is optional.

```
Semaphore Sem("S", 1);  
Sem.Wait();  
// critical section  
Sem.Signal();  
  
Semaphore *Sem;  
Sem = new  
    Semaphore("S", 1);  
Sem->Wait();  
// critical section  
Sem->Signal();
```

# Dining Philosophers: 4 Chairs

```
Semaphore Chairs(4);
Mutex      Chops[5];

class phil::public Thread
{
    public:
        phil(int n, int it);
    private:
        int  Number;
        int  iter;
        void ThreadFunc();
};
```

Count-Down and Lock!

```
Void phil::ThreadFunc()
{
    int i, Left=Number,
        Right=(Number+1)%5;
    Thread::ThreadFunc();
    for (i=0; i<iter; i++) {
        Chairs.Wait();
        Chops[Left].Lock();
        Chops[Right].Lock();
        // Eat
        Chops[Left].Unlock();
        Chops[Right].Unlock();
        Chairs.Signal();
    }
```

# The Smoker Problem: 1/6

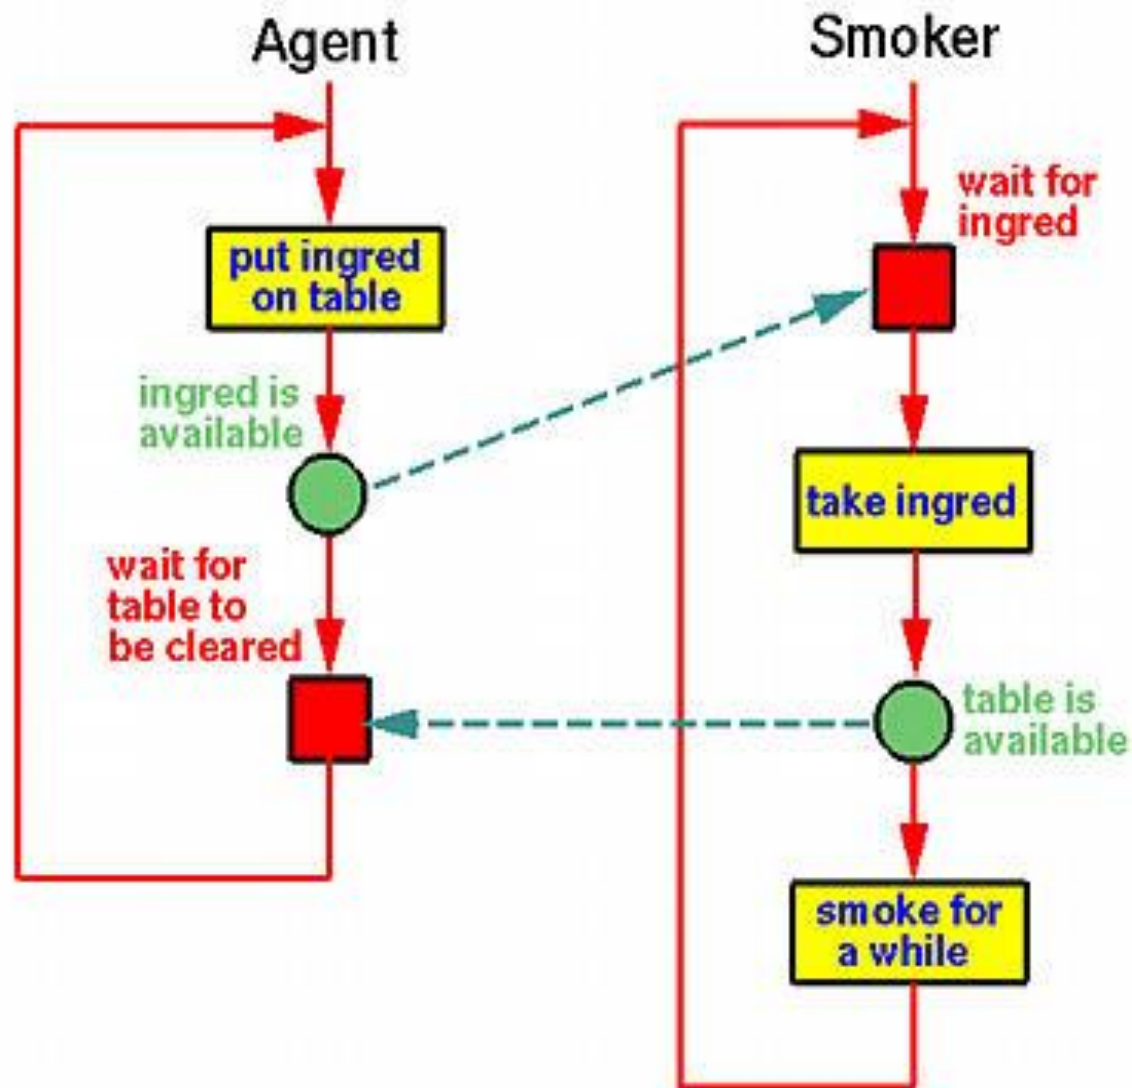
- Three ingredients are needed to make a cigarette: **tobacco**, **paper** and **matches**.
- An agent has an infinite supply of all three.
- Each of the three smokers has an infinite supply of one ingredient only. That is, one of them has tobacco, the second has paper, and the third has matches.
- They share a table.

# The Smokers Problem: 2/6

- The agent adds two randomly selected different ingredients on the table, and notifies the needed smoker.
- A smoker waits until agent's notification. Then, takes the two needed ingredients, makes a cigarette, and smokes for a while.
- This process continues forever.
- **How can we use semaphores to solve this problem?**



# The Smokers Problem: 3/6



# The Smokers Problem: 4/6

- Semaphore `Table` protects the table.
- Three semaphores `Sem[3]` are used, one for each smoker:

<i>Smoker #</i>	<i>Has</i>	<i>Needs</i>	<i>Sem</i>
<b>0</b>	<b>0</b>	<b>1 &amp; 2</b>	<code>Sem[0]</code>
<b>1</b>	<b>1</b>	<b>2 &amp; 0</b>	<code>Sem[1]</code>
<b>2</b>	<b>2</b>	<b>0 &amp; 1</b>	<code>Sem[2]</code>

# The Smokers Problem: 5/6

```
class A::public Thread
{
    private:
        void ThreadFunc();
};
```

*agent thread*

*smoker thread*

```
class Smk::public Thread
{
    public:
        Smk(int n);
    private:
        void ThreadFunc();
        int No;
};
```

*clear the table*

```
Smk::Smk(int n)
{
    No = n;
}
```

```
Void Smk::ThreadFunc()
{
    Thread::ThreadFunc();
    while (1) {
        Sem[No]->Wait();
        Table.Signal();
        // smoker a while
    }
}
```

*waiting for ingredients*

# The Smokers Problem: 6/6

```
void A::ThreadFunc()  
{  
    Thread::ThreadFunc();  
    int Ran;  
    while (1) {  
        Ran = // random #  
              // in [0,2]  
        Sem[Ran]->Signal();  
        Table.Wait();  
    }  
}
```

waiting for the table  
to be cleared



```
void main()  
{  
    Smk *Smoker[3];  
    A Agent;  
  
    Agent.Begin();  
    for (i=0;i<3;i++) {  
        Smoker = new Smk(i);  
        Smoker->Begin();  
    }  
    Agent.Join();  
}
```

**The End**