

Part III

Synchronization

Message Passing

*The most important property of a program is
whether it accomplishes the intention of its user.*

1

Message Passing: 1/3

- When processes/threads are interested with one another, two fundamental requirements must be met: **synchronization** and **communication**.
- **Synchronization** enforces mutual exclusion.
- **Communication** allows information to be passed to other processes/threads.
- Message passing, a form of communication, can be implemented in shared-memory and distributed environment.

Message Passing: 2/3

- Mutex locks, semaphores and monitors are **shared-memory** synchronization mechanisms.
- This means all processes and threads use a piece of shared memory to store and manage mutex locks, semaphores and monitors.
- In a distributed environment, processes and threads run on different computers without a global shared-memory.
- In this case, message passing becomes useful.

Message Passing: 3/3

- Communication links can be established between threads/processes. There are three important issues:
 1. **Naming**: How to refer to each other?
 2. **Synchronization**: Shall we wait when participating a message activity?
 3. **Buffering**: Can messages wait in a communication link?

Naming: Direct Addressing

Symmetric Scheme: 1/3

- **Direct Addressing:** Each process that wants to communicate must explicitly name the other party:
 - ❖ `Send(receiver, message);`
 - ❖ `Receive(sender, message);`
- **With this scheme:**
 - ❖ Exactly one link exists between each pair of communicating processes.
 - ❖ These links may be established for processes that need to communicate before they run.

Naming: Direct Addressing

Asymmetric Scheme: 2/3

- In this scheme, we have
 - ❖ `Send(receiver, message);`
 - ❖ `Receive(id, message);`
- The `Receive()` primitive receives the ID of the sender. Thus, in this scheme, a receiver can receive messages from any process.

Naming: Direct Addressing

Disadvantages: 3/3

- There are disadvantages in the symmetric and asymmetric schemes:
 - ❖ Changing the name/ID of a process may require examining all other process definitions.
 - ❖ Processes must know the IDs of the other parties to start a communication.

Naming: Indirect Addressing

Mailbox: 1/4

- With indirect addressing, messages are sent to and received from **mailboxes**.
- Each mailbox has a unique ID.
- The primitives are
 - ❖ `Send(mailbox-name, message);`
 - ❖ `Receive(mailbox-name, message);`

Naming: Indirect Addressing

Communication Links: 2/4


- **There is a link between two processes only if they share a mailbox.**
- **A link may be shared by multiple processes.**
- **Multiple links may exist between each pair of processes, and each link corresponds to a mailbox.**
- **By decoupling the sender and receiver, indirect addressing provides a greater flexibility in the use of messages.**

Naming: Indirect Addressing

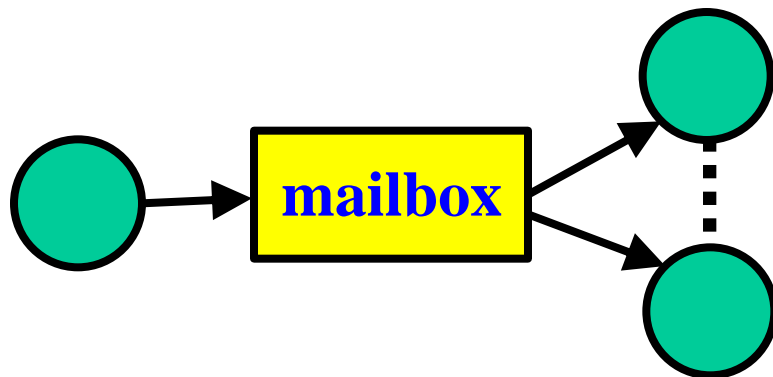
Communication Links: 3/4




one-to-one



many-to-one
mailbox → port



one-to-many



many-to-many

Naming: Indirect Addressing

Communication Links: 4/4

- What if there is only one message in a mailbox and several processes execute `Receive()`? It depends on the following:
 - ❖ If there is only one link between at most two processes, this situation will not happen.
 - ❖ Allow at most one process to receive at a time.
 - ❖ Allow the system to select an arbitrary order.

Synchronization

- The sender and/or receiver may be blocked:
 - ❖ **Blocking Send**: the sender blocks until its message is received
 - ❖ **Nonblocking Send**: the sender sends and resumes its execution immediately
 - ❖ **Blocking Receive**: the receiver blocks until a message is available
 - ❖ **Nonblocking Receive**: the receive receives a message or a null.
- When both send and receive are blocking, we have a **rendezvous** between the sender and receiver.

Synchronous vs. Asynchronous

- **Blocking** and **non-blocking** are known as **synchronous** and **asynchronous**.
 - ❖ If the sender and receiver must **synchronize** their activities, use synchronous communication.
 - ❖ Because of the **uncertainty in the order of events**, asynchronous communication is more difficult to program.
 - ❖ On the other hand, asynchronous algorithms are **general** and **portable**, because they are guaranteed to run correctly on networks with arbitrary timing behavior.

Capacity

- The **capacity** of a link is its **buffer** size:
 - ❖ **Zero Capacity**: Since no message can be waiting in the link, it is **synchronous**. Sender blocks.
 - ❖ **Unbounded Capacity**: Messages can wait in the link. Sender never blocks and the link is **asynchronous**. **The order of messages being received does not have to be FIFO.**
 - ❖ **Bounded Capacity**: Buffered Message Passing. Sender blocks if the buffer is full, and the link is **asynchronous**. Isn't it a bounded buffer if the order is FIFO?

Message Passing in Unix

- Unix systems provide at least two message passing mechanisms: **pipes** and **message queues**.
- A pipe is a generalization of the pipe in `A | B`.
- Function `pipe(int pfd[2])` creates a communication link and returns two file descriptors. Writing to `pfd[1]` puts data in the pipe; reading from `pfd[0]` gets the data out. Data items in a pipe are FIFO just like `A | B`.
- Message queues are mailboxes. Use `msgget()` to obtain a message queue, and use `msgsnd()` and `msgrcv()` to send and receive messages.

Message Passing with **ThreadMentor**

Channels

- A **channel** is a *bi-directional* communication link between two specific threads.
- A channel can be *synchronous* or *asynchronous*.
- Channels use user-defined thread IDs for identification purpose. A use-defined thread ID is a unique non-negative integer selected by the user.
- The user-defined thread ID must be set in the thread constructor:

```
UserDefinedThreadID = 10;
```

Declaring a Channel

- Two classes `SynOneToOneChannel` (blocking send and receive) and `AsynOneToOneChannel` (non-blocking send and receive) are available:

```
SynOneToOneChannel X ("chan-2-3", 15, 3);  
AsynOneToOneChannel Y ("chan-4-5", 15, 3);
```

channel names

user-defined thread IDs

channels `X` and `Y` are built between threads 15 and 3

Sending and Receiving: 1/2

- Use methods `Send()` and `Receive()` to send and receive a message to and from a channel:

```
X.Send(*pointer-to-message, size);
```

```
Y.Receive(*pointer-to-message, size);
```

- Send an `int` message to channel `X`:

```
AsynOneToOneChannel X;
```

```
int Msg;
```

```
X.Send(&Msg, sizeof(int));
```

Sending and Receiving: 2/2

- **Receive a message of four doubles from channel Y:**

```
AsynOneToOneChannel  Y;  
double  Msg[4];  
Y.Receive(Msg, 4*sizeof(double));
```

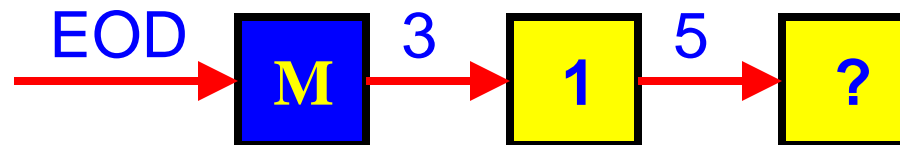
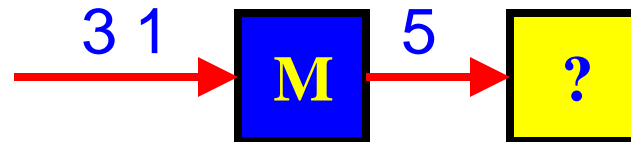
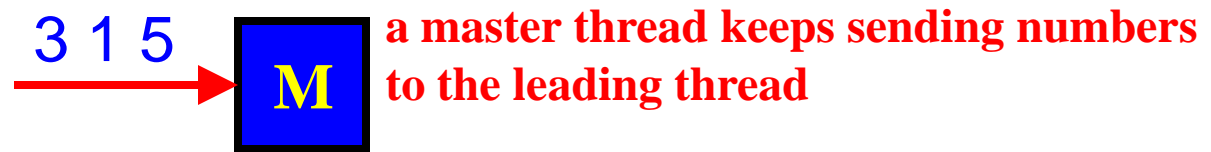
- **Send a message of struct Data to channel Z;**

```
struct Data  
{double x; int y; char z[100];}  
    Msg;  
SynOneToOneChannel  Z;  
Z.Send(&Msg, sizeof(Data));
```

Linear Array Sorting: 1/9

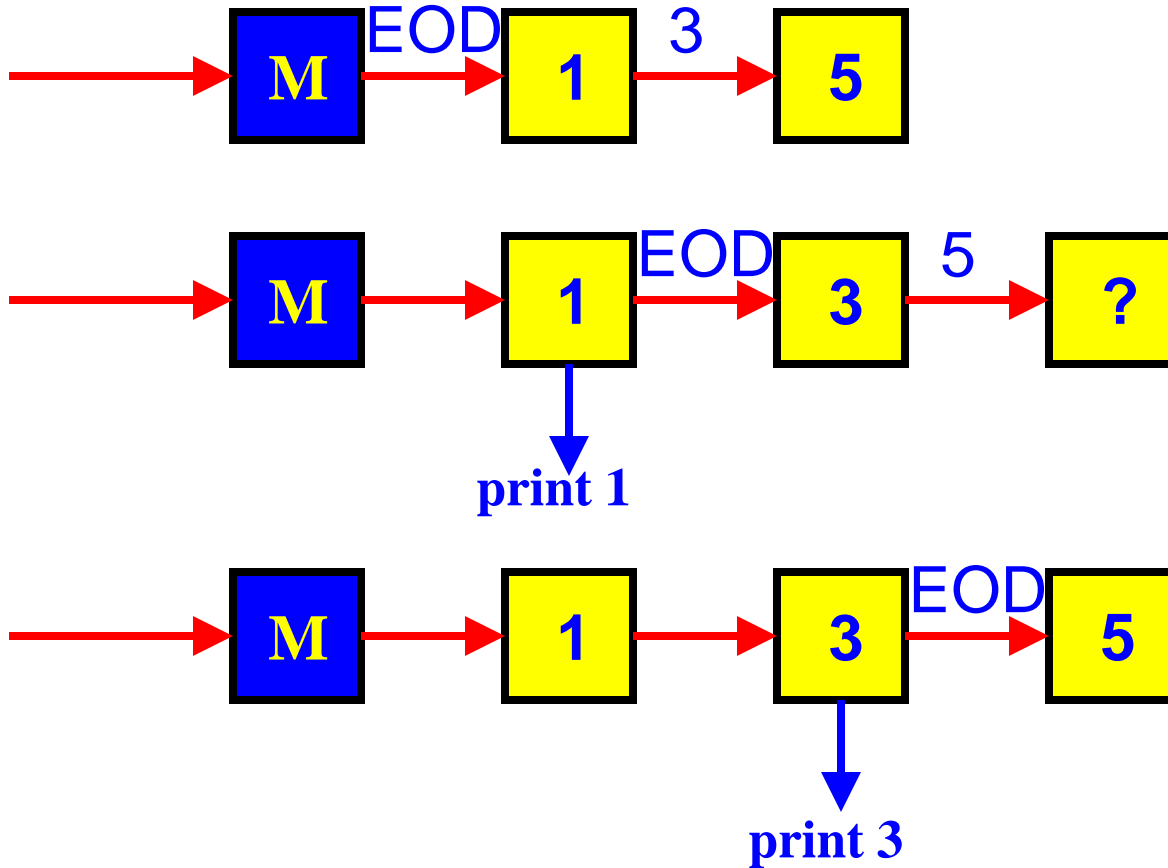
- Each thread has an input channel from its predecessor and an output channel to its successor.
- The first time a thread receives a positive integer, it is memorized as N .
- For a subsequent received positive integer X :
 - ❖ If $X < N$, this thread sends N to its successor and memorizes X as N .
 - ❖ Otherwise, X is sent to its successor.
 - ❖ If there is no successor, this thread creates a thread, builds a channel to it, and sends the number.

Linear Array Sorting: 2/9



sort positive integers with -1 as end-of-data

Linear Array Sorting: 3/9



What type of channels (i.e., sync or async) should be used?

Linear Array Sorting: 4/9

```
const int NOT_DEFINED = -2;
const int END_OF_DATA = -1; // end of input flag
class SortThread : public Thread
{
    public:
        SortThread(int index, int threadID);
        ~SortThread(); // destructor
        SynOneToOneChannel *channel;
    private:
        void ThreadFunc();
        int Index; // index of the sort thread
        int Number; // number memorized
        SortThread *neighbor; // next sort thread
};

class MasterThread : public Thread
{
    public: MasterThread(int threadID);
    private: void ThreadFunc();
};
```

used to construct thread name

class definition

Linear Array Sorting: 5/9

can an async. channel be used here?

```
SortThread::SortThread(int index, int threadID)
{
    Index = index;
    UserDefinedThreadID = threadID;
    neighbor = NULL; // initially no neighbor
    Number = NOT_DEFINED; // no memorized number
    ChannelName = ... // give this channel a name
    channel = new SynOneToOneChannel(ChannelName,
                                     threadID-1, threadID);
}
```

async channel between them

```
SortThread::~~SortThread() // this is a destructor
{ delete channel; }
```

constructor and destructor

```
void SortThread::ThreadFunc()
```

```
{
```

```
    Thread::ThreadFunc();
```

```
    int number, tmpNum;
```

```
    Thread_t self = GetID();
```

```
    while(true) {
```

```
        channel->Receive(&number, sizeof(int)); // receive a number
```

```
        if (number == END_OF_DATA)
```

```
            break;
```

```
        if (Number == NOT_DEFINED)
```

```
            Number = number;
```

```
            // first number. Memorize it
```

```
        else {
```

```
            // other numbers
```

```
            if (number >= Number)
```

```
            // larger than mine
```

```
                tmpNum = number;
```

```
            // save it in temporarily
```

```
            else {
```

```
                tmpNum = Number;
```

```
            // no. save mine in temporarily
```

```
                Number = number;
```

```
            // but, also memorize it
```

```
            }
```

```
        if (neighbor == NULL)
```

```
            // no neighbor? create one
```

```
            neighbor = new SortThread(Index+1, UserDefinedThreadID+1);
```

```
            neighbor->Begin();
```

```
            // run it!
```

```
        }
```

```
        neighbor->channel->Send(&tmpNum, sizeof(int)); // send number
```

```
    }
```

```
} // end of data reached. see next slide
```

GetID() returns the ID of a thread

Linear Array Sorting: 7/9

```
void SortThread::ThreadFunc()
```

SortThread body **2/2**

```
{  
    while (true) {  
        // other stuffs on the previous slide  
        // end of data received  
    }  
    if (neighbor != NULL) { // if I am not the last one  
                            // I should pass the EOD  
        neighbor->channel->Send(&number, sizeof(int));  
        neighbor->Join();    // wait for neighbor to complete  
    }  
    Exit();  
}
```

MasterThread constructor and body

```
MasterThread::MasterThread(int threadID)
{
    UserDefinedThreadID = threadID;
    ThreadName = ...; // a thread name
}

void MasterThread::ThreadFunc()
{
    Thread::ThreadFunc();
    int input;
    do {
        cin >> input; // read an integer or END_OF_DATA
        if (input == END_OF_DATA)
            break;
        else
            firstSortThread->channel->Send(&input, sizeof(int));
    } while (input != END_OF_DATA);
    firstSortThread->channel->Send(&input, sizeof(int));
    Exit();
}
```

send **END_OF_DATA** to the first thread

Linear Array Sorting: 9/9

```
SortThread *firstSortThread; // first sorting thread
void main(void)
{
    MasterThread *masterThread;
    firstSortThread = new SortThread(1,2);
    firstSortThread->Begin();
    masterThread = new MasterThread(1);
    masterThread->Begin();
    masterThread->Join();
    firstSortThread->Join();
    Exit();
}
```

threadID 1

threadID 2

sorting thread #1

main program

The End